


Open PowerQuery

Step 1: Replace Value Fe Male to Female



Step 2: Assign Value to ProductPitched Tier (ordinal)

```
if[ProductPitched]="Basic" then 1  
else if [ProductPitched] = "Standard" then 2  
else if [ProductPitched] = "Deluxe" then 3  
else if [ProductPitched] = "Super Deluxe" then 4  
else 5
```

Step 3: Cap Pitch Duration and replace column

```
if [DurationOfPitch] > 100 then 100 else [DurationOfPitch]
```

Step4: Num of Adults Visiting

```
[NumberOfPersonVisiting] - [NumberOfChildrenVisiting]
```

Step5: Income Per Person

```
[MonthlyIncome] / [NumberOfPersonVisiting]
```

Step 6: Pitch Satisfaction Score

```
=[PitchSatisfactionScore]/[[PitchDuration_Cap]]
```

Step7:

One-Hot Encoding: Converted text columns like [Occupation](#), [Gender](#), [MaritalStatus](#), [Designation](#), and [TypeofContact](#) into binary columns (e.g., [Gender_Male](#): 1 or 0).

Using "1, 2, 3" for categories without a natural rank (like Marital Status) misleads the model because it treats the numbers as having mathematical value, not just as labels. If you code Married=1, Single=2, and Divorced=3, the model assumes "Divorced" is numerically greater than "Married" and might even calculate that the "average" of a Married and Divorced person is a Single person (since $(1+3)/2=2\$$). This false hierarchy forces the model to find patterns based on non-existent logic, whereas splitting them into separate True/False columns (One-Hot Encoding) treats each category as distinct and equal, allowing the AI to learn the actual impact of each specific status independently.

structure and logic of the `assignment5training.py` script. The goal of this script is to train a Deep Learning model (using TensorFlow/Keras) to predict whether a customer will purchase a tourism package (`ProdTaken`).

Part 0: The Helper Function (`split_csv_data`)

What it does

This is a custom function designed to physically separate the raw data into two distinct CSV files before any processing begins.

- **Inputs:** It takes a source file, output filenames, and a `fraction` (set to 0.10 or 10%).
- **Logic:**
 1. It loads the full dataset.
 2. It uses `df.sample(frac=...)` to randomly pick 10% of the rows.
 3. It uses `df.drop(...)` to create a dataset of the remaining 90%.
 4. It saves both as new CSV files.
- **Why is this important?** In Machine Learning, **Data Leakage** is a major problem. By saving the "Test Set" to a completely separate file (`class5_test_data.csv`) at the very beginning, we guarantee that the model *never* sees this data during training. It serves as the "Final Exam."

Part 1: Execution of the Split

```
split_csv_data(..., fraction=0.10, random_state=42)
```

- **random_state=42:** This ensures that every time you run the code, it picks the *exact same* random rows. This makes your results reproducible.
- **Outcome:** You now have `class5_test_data.csv` (10% unseen) and `class5_remaining_data.csv` (90% for training).

Part 2: Loading Training Data

```
df = pd.read_csv(remaining_file_name)
```

- **Crucial Step:** The script intentionally loads the **remaining** 90% file, not the original source. This ensures we are only training on the data meant for learning.

Part 3: Preprocessing

Dropping Columns

- **Index**: Identifier columns do not contain patterns helping prediction, so they are removed.
- **ProductPitched**: This is dropped to prevent **Data Leakage**. If we know which product was pitched, it might be too highly correlated with the result, preventing the model from learning *customer* characteristics.

Defining X and y

- **y (Target)**: `ProdTaken` (Did they buy? 0 = No, 1 = Yes).
- **X (Features)**: All other columns.
- **.astype(float)**: Neural networks require mathematical inputs. This converts any boolean (True/False) data into 1.0 and 0.0.

Part 4: Train / Validation Split

```
X_train, X_test, y_train, y_test = train_test_split(..., test_size=0.25, stratify=y)
```

Even though we already set aside a "Final Exam" file (Part 1), we need an internal "Quiz" to evaluate the model while we are building it.

- **test_size=0.25**: 25% of the *remaining* data is set aside for internal validation.
- **stratify=y**: This is vital for imbalanced data. It ensures that if 18% of people bought the product in the total set, exactly 18% of people in the train set and 18% in the test set also bought it. It keeps the proportions balanced.

Part 5: Scaling

```
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)
```

Neural Networks are very sensitive to the size of numbers.

- **StandardScaler**: Transforms data so the mean is 0 and standard deviation is 1.
- **Why?**: Without this, a column like "Income" (values like 50,000) would overpower a column like "Age" (values like 30).
- **joblib.dump**: Saves the math used for scaling so we can apply the *exact same math* to new data later.

Part 6: Class Weights (Currently Commented Out)

If active, `compute_class_weight` calculates how rare the "Buy" (1) event is. It assigns a higher penalty if the model makes a mistake on a "Buy" prediction, forcing the AI to pay more attention to buyers.

Part 7: The Model Architecture

This is a **Deep Neural Network (DNN)** using the Keras Sequential API.

The Layers

The structure is a "Funnel" shape, getting smaller as it goes deeper:

1. **Dense(512)**: 512 neurons. A wide layer to capture complex patterns.
2. **Dense(256) -> Dense(128) -> ... -> Dense(16)**: Gradually compressing the information to find the most essential features.

Key Components

- **activation='relu'**: The mathematical function that allows the network to learn non-linear patterns (decisions that aren't just straight lines).
- **Dropout(0.4) / Dropout(0.2)**:
 - This is a regularization technique.
 - **0.4** means "Randomly turn off 40% of the neurons in this layer during every training step."
 - **Why?** It prevents the model from memorizing the data (Overfitting) and forces it to learn robust general rules.
- **Output Layer (Dense(1, activation='sigmoid'))**:
 - 1 Neuron: Because we want one answer (Probability of buying).
 - **sigmoid**: Smashes the output between 0 and 1. (e.g., 0.85 means 85% chance of buying).

Part 8: Training with Callbacks

```
model.fit(..., epochs=150, validation_split=0.2, ...)
```

- **epochs=150**: The model sees the data 150 times.
- **validation_split=0.2**: During training, it holds back another 20% of the training data to check its performance at the end of every epoch.


Callbacks (The "Smart" Controls)

1. **EarlyStopping**:

- Monitors `accuracy`.
 - If the accuracy doesn't improve for 10 epochs (`patience=10`), it stops training immediately.
 - `restore_best_weights=True`: It reverts the model back to its "best" state, ignoring the last 10 bad epochs.
2. **ReduceLROnPlateau:**
- If the model gets stuck based on what it monitors, it cuts the `learning_rate` in half (`factor=0.5`).
 - This allows the model to make smaller, more precise adjustments to try and find a better solution.

Part 9 & 10: Evaluation and Saving

- `model.evaluate`: Checks the final accuracy on the `X_test` data (the 25% internal split).
- `classification_report`: Shows Precision (Accuracy of positive predictions) and Recall (How many of the actual positives were found).
- `model.save`: Saves the trained brain of the AI to a `.keras` file so it can be loaded later without retraining.



INFERENCE RESULTS

Accuracy: 0.8972

AUC Score: 0.9053

Classification Report:

| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

| | | | | |
|---|------|------|------|-----|
| 0 | 0.94 | 0.93 | 0.94 | 256 |
| 1 | 0.74 | 0.77 | 0.75 | 65 |

| | | |
|----------|------|-----|
| accuracy | 0.90 | 321 |
|----------|------|-----|

| | | | | |
|--------------|------|------|------|-----|
| macro avg | 0.84 | 0.85 | 0.84 | 321 |
| weighted avg | 0.90 | 0.90 | 0.90 | 321 |

Confusion Matrix:

`[[238 18]`

`[15 50]]`